




Getting Started with Docker

Containers, Images, Networking,
Volumes, Compose and Security

Docker

Podman

Networking

Volumes

Security

Salvatore D'Angelo

code4projects.com

2026

Contents

1	Preface	5
1.1	Who This Book Is For	5
1.2	How This Book Is Structured	5
1.3	A Note on Docker vs Podman	6
1.4	Source Code and Updates	6
2	Getting Started with Docker and Podman	7
2.1	Introduction	7
2.2	Why Docker?	7
2.3	Docker vs Podman: Which One Should You Use?	8
2.4	Installing Docker or Podman	9
2.5	Core Concepts: Images and Containers	9
2.6	Essential Commands	10
2.7	Your First Containerized Application	12
2.8	How This Series Is Structured	13
2.9	Conclusion	13
3	Containers vs Virtual Machines	14
3.1	Introduction	14
3.2	Virtual Machines	14
3.3	Containers	16
3.4	The Analogy	17
3.5	Containers and VMs: not enemies, complementaries	18
3.6	Conclusion	18
4	Dockerfile and Building Custom Images	19
4.1	Introduction	19
4.2	What is a Dockerfile?	19
4.3	Dockerfile Instructions	19
4.4	Building a Custom Nginx Image	21
4.5	Understanding the Layer Cache	22
4.6	Tagging and Versioning Images	23
4.7	Cleaning Up	24
4.8	Conclusion	24
5	How Docker Networking Works	25
5.1	Introduction	25
5.2	Networking Overview in Docker	25
5.3	Docker Networking Commands	27
5.4	Practical Example: Frontend and Backend on the Same Network	27
5.5	Conclusion	29
6	How Docker Volumes Work	30
6.1	Introduction	30
6.2	The Upgrade Problem	30
6.3	Docker Storage Types	31
6.4	The <code>--mount</code> Syntax	32
6.5	Docker Volumes Cheat Sheet	33

6.6	Practical Example: Nginx with Persistent Content	33
6.7	Volumes vs Bind Mounts: When to Use Which	34
6.8	Conclusion	35
7	How Docker Compose Works	36
7.1	Introduction	36
7.2	What is Docker Compose?	36
7.3	The docker-compose.yml File	37
7.4	Practical Example: Frontend and Backend with Compose	37
7.5	Docker Compose Commands	39
7.6	Environment Variables	40
7.7	Starting the Application	41
7.8	Conclusion	41
8	Docker Security Best Practices	43
8.1	Introduction	43
8.2	1. Never Run as Root	43
8.3	2. Use a Read-Only Filesystem	45
8.4	3. Drop Linux Capabilities	45
8.5	4. Never Put Secrets in Environment Variables (or Images)	46
8.6	5. Use Minimal Base Images	46
8.7	6. Scan Images for Known Vulnerabilities	47
8.8	7. Keep Images Up to Date	47
8.9	Putting It All Together	48
8.10	Conclusion	49
9	Glossary	50
10	What's Next?	52
10.1	Where to Go From Here	52
10.2	Stay in Touch	53

To every developer who has ever typed
`docker run hello-world`
and wondered what actually just happened.

1 Preface

I started working with Docker in 2018 when containers were already gaining momentum but most tutorials either assumed too much prior knowledge or buried the practical steps under layers of theory.

This book is what I wish I had back then.

It is a no-nonsense, hands-on introduction to Docker and Podman. Every chapter introduces one concept and immediately applies it with a concrete example. By the time you finish the last chapter, you will have built, networked, persisted, orchestrated, and secured a real multi-container application — step by step, command by command.

1.1 Who This Book Is For

This book is for developers who:

- Have heard of Docker but have never used it in practice
- Want to understand the fundamentals before jumping into Kubernetes
- Work on a corporate machine where Docker Desktop requires a paid license and need a free alternative (Podman)
- Want a single, coherent reference instead of scattered blog posts

No prior knowledge of containers or virtualization is assumed. Basic familiarity with the command line is all you need.

1.2 How This Book Is Structured

Each chapter of this book corresponds to a foundational Docker concept, introduced in a deliberate order so that each one builds on the previous:

1. **Getting Started with Docker and Podman** — first contact, core concepts, essential commands, your first running container.
2. **Containers vs Virtual Machines** — how containers actually work under the hood, what Linux namespaces and cgroups are, and when to use VMs instead.
3. **Dockerfile and Building Custom Images** — write your own Dockerfile, understand the layer cache, build and tag images.
4. **How Docker Networking Works** — connect containers so they can talk to each other, automatic DNS resolution, network drivers.

5. **How Docker Volumes Work** — separate data from code, persist state across container restarts and replacements.
6. **How Docker Compose Works** — manage a multi-container application with a single YAML file and a single command.
7. **Docker Security Best Practices** — run containers safely: non-root users, read-only filesystems, capability limits, vulnerability scanning.

1.3 A Note on Docker vs Podman

All commands in this book are shown with the `docker` prefix. If you are using Podman, every command works identically — simply replace `docker` with `podman`, or set the alias `alias docker=podman` once and never think about it again.

1.4 Source Code and Updates

The examples in this book are maintained on GitHub at github.com/sasadangelo/docker-tutorials.

For corrections, suggestions, or feedback, write to sasadangelo@gmail.com or visit code4projects.com.

2 Getting Started with Docker and Podman



Figure 1: Getting Started with Docker and Podman

2.1 Introduction

This is the first article of the **Getting Started with Docker** series. By the end of this article you will understand what Docker is, what problems it solves, and you will have your first containerized web server running on your machine in minutes — without writing a single line of Dockerfile.

You should read this article if:

- You have heard of Docker but never used it
- You want a practical, no-nonsense introduction that gets you running immediately
- You are on a machine where Docker Desktop is not available and you need an alternative (Podman)

All the commands shown here work identically with both **Docker** and **Podman**. Where there are differences, they are explicitly noted.

2.2 Why Docker?

Before diving into commands, it is worth spending a minute on the problem Docker actually solves.

Imagine you are a developer. You write an application on your laptop, it works perfectly. Then you hand it off to a colleague or deploy it to a server — and it breaks. Different OS version, different library installed, different path. This is the classic “**it works on my machine**” problem, and it has tormented developers for decades.

Docker solves this by packaging your application together with everything it needs to run — runtime, libraries, configuration — into a single portable unit called a **container**. That container runs identically on your laptop, your colleague’s machine, and the production server.

This matters not just for individual developers. Modern applications are built as dozens (sometimes hundreds) of independent services — **microservices** — each with its own dependencies. Docker makes it possible to deploy, update, and scale each service independently, without interference.

2.3 Docker vs Podman: Which One Should You Use?

Both Docker and Podman are tools for running containers. For everything covered in this series, they are interchangeable — the CLI commands are identical.

The key differences:

	Docker	Podman
Architecture	Requires a background daemon (dockerd)	Daemonless — runs containers directly
Privileges	Daemon runs as root	Rootless by default — more secure
License	Docker Desktop requires a paid license for enterprise use	Fully open source (Apache 2.0)
Installation on Mac	Docker Desktop	<code>brew install podman + podman machine init</code>
Compose support	<code>docker compose</code> (v2 plugin)	<code>podman compose</code> (from Podman 4.7+)
OCI compatibility	Yes	Yes — same image format

In short: if Docker Desktop is available and licensed for your use, use Docker. If you are on a corporate machine where Docker Desktop requires a commercial license, or if you prefer an open source rootless alternative, use Podman. The rest of this series works exactly the same either way.

Throughout this series, all commands are shown with `docker`. If you are using Podman, simply replace `docker` with `podman` — or set the alias `alias docker=podman` once and forget about it.

2.4 Installing Docker or Podman

2.4.1 Docker

On Linux, install Docker Engine following the [official documentation](#). On Mac and Windows, Docker Desktop is the easiest option if your license allows it: download it from docs.docker.com/desktop.

2.4.2 Podman (Mac)

On macOS, Podman uses a lightweight Linux virtual machine behind the scenes — the same approach Docker Desktop uses. Install and initialise it with three commands:

```
brew install podman
podman machine init
podman machine start
```

After this, podman is ready and behaves exactly like docker.

2.5 Core Concepts: Images and Containers

Before running anything, two concepts are essential.

2.5.1 Docker Image

A **Docker image** is a read-only, layered package that contains everything needed to run a specific application: the OS filesystem, runtime, libraries, configuration files, and the application itself.

Images are stored in registries. The most popular is [Docker Hub](#), which hosts thousands of official and community images. When you run a container, Docker (or Podman) downloads the image from the registry if it is not already present on your machine.

Images are built in **layers**. Each instruction that modifies the filesystem adds a layer on top of the previous ones. Layers are cached and shared between images — this is what keeps images small and builds fast.

2.5.2 Docker Container

A **container** is a running instance of an image. You can think of the image as a class and the container as an object instantiated from it. Multiple containers can run from the same image simultaneously, each fully isolated from the others.

Containers share the host OS kernel but are isolated from each other using Linux kernel features: **namespaces** (for process, network, filesystem isolation) and **cgroups** (for resource limits like CPU and memory).

This is the fundamental difference from a Virtual Machine: a VM emulates an entire hardware stack and runs a separate OS kernel. A container shares the host kernel and is much lighter as a result.

We will explore this topic in depth in the [next article](#).

2.6 Essential Commands

Here is the minimal set of commands you need to start working with Docker. These are the ones you will use every day.

2.6.1 List local images

```
docker image ls
```

At the beginning, the list is empty.

2.6.2 Pull an image from Docker Hub

```
docker pull nginx:alpine
```

This downloads the official Nginx image based on Alpine Linux — a very small (~40 MB) but fully functional web server image.

2.6.3 Run a container

```
docker run -d -p 8080:80 --name my-nginx nginx:alpine
```

Breaking down the options:

- `-d` — runs the container in detached mode (in the background, shell does not hang)
- `-p 8080:80` — maps port 8080 on your host to port 80 inside the container

- `--name my-nginx` — gives the container a friendly name instead of a random one
- `nginx:alpine` — the image to use

Open your browser at `http://localhost:8080` and you will see the Nginx welcome page. You just ran a web server without installing anything on your machine.

2.6.4 List running containers

```
docker container ls
```

To see all containers including stopped ones:

```
docker container ls -a
```

2.6.5 View container logs

```
docker logs my-nginx
```

2.6.6 Execute a command inside a running container

```
docker exec -it my-nginx /bin/sh
```

The `-it` flags attach an interactive terminal. Note: Alpine-based images use `sh` instead of `bash`.

2.6.7 Stop and remove a container

```
docker stop my-nginx  
docker rm my-nginx
```

Or in one command:

```
docker rm -f my-nginx
```

2.6.8 Remove an image

```
docker image rm nginx:alpine
```

2.6.9 Clean up everything unused

```
docker system prune
```

This removes all stopped containers, dangling images, and unused networks. Add `-a` to also remove unused images that are not referenced by any container.

2.7 Your First Containerized Application

Let us put it all together with a practical example. You will serve a custom HTML page through Nginx — without installing Nginx, without a Dockerfile, using only the official image and a single `docker run` command.

Create a file called `index.html` anywhere on your machine:

```
Hello from Docker

Hello, World!

This page is served by Nginx running inside a container.
```

Now run Nginx and mount that file into the container:

```
docker run -d \
  -p 8080:80 \
  --name hello-nginx \
  -v $(pwd)/index.html:/usr/share/nginx/html/index.html:ro \
  nginx:alpine
```

The `-v` option **bind mounts** your local file into the container's web root. The `:ro` suffix makes the mount read-only inside the container (a good practice).

Open `http://localhost:8080` in your browser:

Hello, World! This page is served by Nginx running inside a container.

You served a custom page without installing a web server, without writing a Dockerfile, and without touching your system configuration. When you are done:

```
docker rm -f hello-nginx
```

Your machine is exactly as it was before. Nothing installed, nothing left behind.

2.8 How This Series Is Structured

This first article introduced the core concepts and got you running immediately. In the upcoming articles we will progressively build on this foundation:

- **Containers vs Virtual Machines** — deep dive into isolation mechanisms, namespaces, cgroups, and the difference from hypervisor-based virtualization
- **Dockerfile & Building Custom Images** — write your own `Dockerfile`, build and tag a custom image, understand the layer cache
- **How Docker Networking Works** — user-defined bridge networks, automatic DNS resolution between containers, network isolation
- **How Docker Volumes Work** — named volumes vs bind mounts, persisting data across container restarts
- **How Docker Compose Works** — manage multi-container applications with a single `docker compose up`
- **Docker Security Best Practices** — run containers as non-root, drop capabilities, scan images for vulnerabilities

Each article builds on the previous one around the same practical Nginx application, adding one concept at a time.

2.9 Conclusion

In this article we covered:

- The problem Docker solves and why containers matter
- The difference between Docker and Podman, and when to use each
- The two fundamental concepts: **images** and **containers**
- The essential CLI commands for daily use
- Running a custom web page with Nginx using only the official image

The **next article** goes deeper into the conceptual foundations: what actually makes a container different from a virtual machine, and what Linux kernel mechanisms make it all work.

If you enjoyed this article, don't forget to **give it a clap** , **share it with your friends** , and **follow me for more tips and tutorials on software development** . Your support helps me create more content like this — thank you!

3 Containers vs Virtual Machines




Figure 2: Containers vs Virtual Machines

3.1 Introduction

This is the second article of the **Getting Started with Docker** series. In the [previous article](#) we introduced Docker and ran our first container. Before going further, it is worth spending time on a fundamental question: **what is the real difference between a container and a virtual machine?**

When I started working with Docker this was not easy to grasp. The two technologies look similar from the outside — both give you an isolated environment to run software — but they work in completely different ways under the hood.

Virtualization can be achieved in two ways:

- **Full virtualization** — a hypervisor simulates enough hardware to run a complete guest operating system. This is the Virtual Machine approach.
- **Operating system-level virtualization** — the OS kernel itself is used to create isolated environments that share the same kernel. This is the Container approach.

3.2 Virtual Machines

Historically, as server processing power grew, applications running on bare metal couldn't exploit the abundance of available resources. **Virtual Machines** were born to solve this: by

running a hypervisor on top of physical hardware, you could carve up one server into many independent guest systems.

A **hypervisor** (also called Virtual Machine Monitor) is the software layer that sits between the hardware and the virtual machines. It can be type 1 (bare metal, like VMware ESXi or KVM) or type 2 (hosted, like VirtualBox or Parallels).

Each VM includes:

- A full copy of a guest operating system
- Virtual hardware (CPU, RAM, disk, network)
- All the application dependencies

The system running the VMs is called the **host**, the systems running inside are called **guests**.

3.2.1 Pros and cons of VMs

Pros

Strong isolation – each VM has its own kernel.

Run different operating systems on the same host (Windows on Linux, etc.).

Create and destroy on demand in minutes.

Snapshot a VM at any point in time.

Foundation of the IaaS cloud model (AWS EC2, Azure VMs, etc.).

Cons

Heavy – each VM carries a full OS (GBs of disk, hundreds of MB of RAM).

Slow to boot – minutes, not seconds.

Hypervisor adds overhead to application execution.

Managing a large fleet of VMs is operationally expensive.

The rise of VMs gave birth to the **Infrastructure as a Service (IaaS)** market. Companies like Amazon, Microsoft, and IBM allow you to rent virtual machines on demand. A startup today

does not need to buy physical hardware — it registers an account with a cloud provider, orders a VM with the required specs, and is ready in minutes.




Figure 3: VM vs Container stack — VM has a full Guest OS per application, containers share the host kernel

3.3 Containers

A **container** is an isolated environment created using features of the Linux kernel itself — not a separate hypervisor. Multiple containers run on the same host, all sharing the same OS kernel, but each seeing its own isolated filesystem, processes, network, and users.

The two kernel mechanisms that make this possible are:

- **Namespaces** — give each container an isolated view of the system: its own process tree (PID namespace), network interfaces (net namespace), filesystem (mount namespace), users (user namespace), and more.
- **cgroups** (control groups) — limit and account for resource usage (CPU, memory, disk I/O) per container.

Together, namespaces and cgroups create the illusion of a separate machine without the overhead of emulating hardware or running a second kernel.

Important note for Mac and Windows users: because containers rely on Linux kernel features, Docker and Podman on macOS and Windows run a lightweight Linux VM behind the scenes (using Apple's Virtualization Frame-

work on Mac, and WSL2 on Windows). Your containers still run in that Linux environment — you just don't see it.

3.3.1 Pros and cons of containers

Pros

Lightweight — a container image is tens or hundreds of MB, not GBs.

Fast — containers start in milliseconds to seconds.

No overhead from emulating hardware.

Portable — same image runs on your laptop, CI server, and production.

Solves the "it works on my machine" problem.

Natural fit for microservices architecture.

Cons

Weaker isolation than VMs — all containers share the host kernel.

Cannot run a different OS kernel (e.g., no Windows containers on Linux).

Managing many containers requires an orchestrator (Kubernetes, Swarm).

Security requires care — a misconfigured container can affect the host.

3.4 The Analogy

Mike Coleman on the [Docker official blog](#) provides a memorable analogy:

Houses (VMs) are fully self-contained and offer protection from unwanted guests. They each possess their own infrastructure — plumbing, heating, electrical. Apartments (containers) also offer protection from unwanted guests, but they are built around shared infrastructure. The apartment building (Docker Host) shares plumbing, heating, electrical. Apartments are offered in all kinds of sizes — from studio to penthouse. You only rent exactly what you need.

3.5 Containers and VMs: not enemies, complementaries

In practice, containers and VMs are often used **together**, not in opposition. The most common production setup is:

1. A cloud provider gives you a VM (an EC2 instance, an Azure VM, etc.)
2. Docker or a container runtime runs inside that VM
3. Your application runs as containers inside Docker

VMs provide the strong isolation boundary at the infrastructure level. Containers provide the lightweight, fast, portable packaging for applications within that boundary.

3.6 Conclusion

In this article we covered:

- The difference between full virtualization (VMs) and OS-level virtualization (containers)
- How **namespaces** and **cgroups** implement container isolation
- Pros and cons of each technology
- Why containers and VMs are complementary, not competing

The **next article** focuses on **Dockerfiles**: how to write your own image definition and build a custom container image from scratch.

If you enjoyed this article, don't forget to **give it a clap** , **share it with your friends** , and **follow me for more tips and tutorials on software development** . Your support helps me create more content like this — thank you!

4 Dockerfile and Building Custom Images




Figure 4: Dockerfile and Building Custom Images

4.1 Introduction

This is the third article of the **Getting Started with Docker** series. In the [previous article](#) we explored the differences between containers and virtual machines. In the [first article](#) we ran Nginx directly from the official image using a bind mount to serve a custom page.

That approach works well for quick experiments, but in a real project you want a **reproducible, self-contained image** that already includes your content and configuration — one that anyone on your team (or any CI/CD pipeline) can build and run without manual steps.

This is where **Dockerfiles** come in.

4.2 What is a Dockerfile?

A **Dockerfile** is a plain-text file containing a sequence of instructions that tell Docker how to build an image. Each instruction creates a new **layer** on top of the previous one. When you rebuild an image, Docker only rebuilds the layers that changed — everything before the first change is served from the **layer cache**, making subsequent builds very fast.

4.3 Dockerfile Instructions

Here are the most important instructions you will use in every Dockerfile.

4.3.1 FROM

```
FROM nginx:alpine
```

Every Dockerfile must start with FROM. It specifies the **base image** — the starting point. In this series we always start from `nginx:alpine`, a minimal image that already contains a working Nginx web server.

4.3.2 COPY

```
COPY ./html /usr/share/nginx/html
```

Copies files or directories from your build context (the folder where you run `docker build`) into the image. This is the main instruction for adding your application code or configuration.

4.3.3 ADD

ADD is similar to COPY but with two extra powers: it can automatically unpack compressed archives (`.tar.gz`, etc.) and it can fetch files from remote URLs. For simply copying local files, **prefer COPY** — it is explicit and predictable.

4.3.4 RUN

```
RUN apk add --no-cache curl
```

Executes a command **during the build** and commits the result as a new layer. Use it to install packages, compile code, or run any setup step. Each RUN instruction is a layer, so chain related commands with `&&` to minimize layer count.

4.3.5 ENV

```
ENV APP_ENV=production
```

Sets an environment variable that is available both during the build and at container runtime.

4.3.6 EXPOSE

```
EXPOSE 80
```

Documents which port the container listens on. It does **not** actually publish the port — that happens at runtime with `-p`. Think of it as metadata for the image user.

4.3.7 CMD

```
CMD ["nginx", "-g", "daemon off;"]
```

Specifies the default command to run when a container starts. It can be overridden at runtime. Use the **exec form** (JSON array) rather than the shell form to avoid wrapping the process in a shell — this ensures signals like `SIGTERM` are delivered directly to your process.

4.3.8 ENTRYPOINT

Similar to `CMD` but not easily overridden. Use `ENTRYPOINT` when the container has a fixed main executable, and `CMD` to provide default arguments to it.

4.3.9 USER

```
USER nginx
```

Switches to a non-root user for all subsequent instructions and at runtime. This is a **security best practice** — we will cover it in detail in the security article.

4.4 Building a Custom Nginx Image

Let us build a custom image that already includes our HTML page, so anyone can run it with a single `docker run` — no bind mounts, no extra files needed.

Create a project directory:

```
mkdir my-nginx && cd my-nginx
```

Create the HTML content:

```
mkdir html
```

```
cat > html/index.html
```

```
My Custom Nginx
```

```
    Hello from my custom image!
```

```
    This page is baked into the Docker image.
```

EOF

Create the Dockerfile:

```
FROM nginx:alpine

# Copy our HTML content into the image
COPY html /usr/share/nginx/html

# Document which port Nginx listens on
EXPOSE 80
```

Build the image:

```
docker build -t my-nginx:1.0 .
```

Breaking down the command:

- `-t my-nginx:1.0` — **tag** the image with a name and version
- `.` — the **build context** (current directory); Docker sends all files here to the build engine

List your images to confirm it was created:

```
docker image ls
```

Run a container from your custom image:

```
docker run -d -p 8080:80 --name my-app my-nginx:1.0
```

Open `http://localhost:8080` — you will see your custom page, served directly from the image without any bind mounts.

4.5 Understanding the Layer Cache

Run `docker build` a second time without changing anything:

```
docker build -t my-nginx:1.0 .
```

You will see output like:

```
=> CACHED [1/2] FROM docker.io/library/nginx:alpine
=> CACHED [2/2] COPY html /usr/share/nginx/html
```



Figure 5: Dockerfile instructions and image layers

Both steps are served from cache — the build is instant. Now edit `html/index.html` and rebuild:

```
docker build -t my-nginx:1.1 .
```

Only the COPY layer is rebuilt. The FROM layer (downloading the base image) is still cached.

Rule of thumb: order your Dockerfile instructions from least to most frequently changing. Put RUN commands that install packages near the top, and COPY of your application code near the bottom. This maximises cache reuse and keeps builds fast.

4.6 Tagging and Versioning Images

Good image tags tell you exactly what version is inside:

```
# Tag with a version number
docker build -t my-nginx:1.0 .

# Tag the same image as latest
docker tag my-nginx:1.0 my-nginx:latest
```

Always use explicit version tags in production. The latest tag is convenient locally but can cause surprises in automated pipelines because it changes silently.

4.7 Cleaning Up

Remove the running container:

```
docker rm -f my-app
```

Remove the image:

```
docker image rm my-nginx:1.0 my-nginx:1.1
```

Remove all unused images and build cache:

```
docker system prune -a
```

4.8 Conclusion

In this article we covered:

- The purpose of a **Dockerfile** and how it defines an image
- The most important instructions: FROM, COPY, RUN, ENV, EXPOSE, CMD, ENTRYPOINT, USER
- How to **build**, **tag**, and **run** a custom Nginx image
- How the **layer cache** works and how to structure a Dockerfile to maximise it

The **next article** introduces Docker networking: how to connect containers together so they can communicate over a private network.

If you enjoyed this article, don't forget to **give it a clap** , **share it with your friends** , and **follow me for more tips and tutorials on software development** . Your support helps me create more content like this — thank you!

5 How Docker Networking Works




Figure 6: How Docker Networking Works

5.1 Introduction

This is the fourth article of the **Getting Started with Docker** series. In the [previous article](#) we built a custom Nginx image. So far, each container has been an isolated island. In a real application you often have multiple containers that need to talk to each other — a frontend container, a backend container, a database.

This is where Docker networking comes in.

5.2 Networking Overview in Docker

Docker has a pluggable networking system. The networking plugins are called **drivers**. Docker provides these drivers out of the box:

1. **bridge** — the default; creates a private internal network on the host
2. **host** — the container shares the host's network stack directly
3. **overlay** — connects containers running on different hosts (used in Swarm mode)
4. **macvlan** — assigns a real MAC address to a container, making it appear as a physical device on the network
5. **ipvlan** — similar to macvlan but operates at layer 3; useful when the upstream switch does not allow multiple MACs per port

6. **none** — networking is completely disabled

For this series we focus on **bridge**, which is what you will use for local development and single-host deployments.

5.2.1 The bridge driver

The **bridge** driver creates a private network on the host. Containers attached to the same bridge can communicate with each other. External access is granted by publishing ports with `-p`.

Behind the scenes, Docker creates a Linux bridge interface, assigns IP addresses from a private subnet, and manages iptables rules to route traffic in and out.

There are two kinds of bridge networks:

- **default bridge** (bridge) — the one Docker uses automatically when you run a container without specifying a network. Containers can reach each other only by IP address, not by name.
- **user-defined bridge** — created with `docker network create`. Containers on the same user-defined bridge can **resolve each other by container name** via Docker's built-in DNS. This is the right choice for any multi-container setup.

5.2.2 The host driver

With the **host** driver the container skips Docker's network layer entirely and shares the host's network interfaces. Port mapping (`-p`) is not needed — if the container listens on port 80, it is directly reachable on the host's port 80.

Useful for performance-sensitive workloads or tools that need to monitor host network traffic. Note: `--network host` does not work on Docker Desktop for Mac or Windows — it only works natively on Linux.

5.2.3 The overlay driver

The **overlay** driver enables containers on different physical hosts to communicate as if they were on the same network. It is the networking backbone for **Docker Swarm** clusters. We will cover this when we discuss Swarm in a future article.

5.2.4 The macvlan driver

The **macvlan** driver assigns a MAC address to a container. This makes it appear as a physical device directly on the network segment. It is the right choice for legacy applications that expect

to be directly connected to the physical network rather than routed through Docker's network stack.

5.3 Docker Networking Commands

Create a user-defined bridge network:

```
docker network create my-network
```

To specify a subnet and gateway explicitly:

```
docker network create \
  -d bridge \
  --subnet=10.10.1.0/24 \
  --gateway=10.10.1.1 \
  my-network
```

List all networks:

```
docker network ls
```

Inspect a network (shows connected containers, subnet, gateway):

```
docker network inspect my-network
```

Remove a network:

```
docker network rm my-network
```

Connect a running container to a network:

```
docker network connect my-network my-container
```

5.4 Practical Example: Frontend and Backend on the Same Network

Let us build a concrete example. We will run two containers on the same user-defined bridge network and have them communicate by name — no IP addresses needed.

5.4.1 Create the network

```
docker network create app-network
```




Figure 7: Docker networking — frontend and backend on the same bridge

5.4.2 Start the “backend” container

For simplicity we use another Nginx instance as a mock backend serving a JSON response.

```
mkdir backend
cat > backend/index.html frontend/index.html

Frontend

Frontend container
Backend is reachable at http://backend

EOF

docker run -d \
  --name frontend \
  --network app-network \
  -p 8080:80 \
  -v $(pwd)/frontend:/usr/share/nginx/html:ro \
  nginx:alpine
```

5.4.3 Verify containers can resolve each other by name

```
docker exec frontend ping -c 3 backend
```

Because both containers are on the same user-defined bridge network, Docker's built-in DNS resolves backend to the correct container IP automatically. No need to hardcode IP addresses anywhere.

Open `http://localhost:8080` in your browser to see the frontend. The backend is reachable internally at `http://backend` from within the frontend container.

5.4.4 Clean up

```
docker rm -f frontend backend
docker network rm app-network
```

5.5 Conclusion

In this article we covered:

- Docker's five network drivers: **bridge**, **host**, **overlay**, **macvlan**, **ipvlan**, and **none**
- The difference between the **default bridge** and **user-defined bridge** networks
- How **automatic DNS resolution** works on user-defined bridges (containers talk by name, not IP)
- Essential networking commands: `create`, `ls`, `inspect`, `rm`, `connect`
- A practical example with frontend and backend containers communicating by name

The [next article](#) covers **Docker volumes**: how to persist data outside a container so it survives restarts and replacements.

If you enjoyed this article, don't forget to **give it a clap** , **share it with your friends** , and **follow me for more tips and tutorials on software development** . Your support helps me create more content like this — thank you!

6 How Docker Volumes Work




Figure 8: How Docker Volumes Work

6.1 Introduction

This is the fifth article of the **Getting Started with Docker** series. In the [previous article](#) we connected containers over a private network. Now we need to address another fundamental question: **what happens to your data when a container is stopped or replaced?**

By default, a container's filesystem is **ephemeral** — everything written inside it disappears when the container is removed. This is fine for stateless services like a web server, but it is a problem as soon as you have anything that needs to survive: logs, configuration, user data, database files.

Docker solves this with **volumes**.

6.2 The Upgrade Problem

One of Docker's main benefits is that upgrading an application is as simple as replacing the old container with a new one. But if your data is stored inside the container, replacing it means losing the data.

The solution is to **separate the application binaries from the data**. Store the binaries in the image and the data in a volume that lives outside the container. When you replace the container, the volume stays where it is and the new container picks it up.

6.3 Docker Storage Types

Docker provides three options for persisting data from a container to the host filesystem.

6.3.1 Volumes

Volumes are managed entirely by Docker. They are stored in a directory on the host that Docker controls (typically `/var/lib/docker/volumes/` on Linux) and you interact with them through Docker CLI commands — not by poking around the filesystem directly.

Volumes are the **recommended way to persist data** in Docker. From the [official documentation](#), some of the reasons:

- Easier to back up and migrate than bind mounts
- Manageable through the Docker CLI and API
- Work on both Linux and Windows containers
- Can be safely shared among multiple containers
- Support remote storage drivers (NFS, cloud block storage, etc.)
- New volumes can be pre-populated with data from a container

6.3.2 Bind Mounts

Bind mounts map a specific directory or file on the **host filesystem** into the container. Unlike volumes, bind mounts are not managed by Docker — any process on the host can read and write them.

Bind mounts are ideal when:

- You want to share source code between your editor on the host and the running container (live reload during development)
- You need the storage to be on a specific filesystem type
- You need direct access to the files from the host without going through Docker

We already used a bind mount in the first article: `-v $(pwd)/index.html:/usr/share/nginx/html/index.ht`

6.3.3 Tmpfs Mounts

Tmpfs mounts store data in the host's memory, not on disk. The data is never written to disk and disappears when the container stops. They are useful for sensitive data (secrets, tokens) that should never touch disk, or for temporary scratch space that requires very fast I/O.

Tmpfs mounts are only available on Linux hosts.




Figure 9: Docker storage types: volumes, bind mounts, tmpfs

6.4 The `--mount` Syntax

Docker has two ways to attach storage to a container: the older `-v` flag and the newer `--mount` flag. The `--mount` syntax is more explicit and is the one recommended by the official documentation today.

Comparison:

```
# Old syntax (-v)
docker run -v my-volume:/app/data nginx:alpine

# New syntax (--mount) — equivalent
docker run --mount type=volume,source=my-volume,target=/app/data nginx:alpine
```

The `--mount` form accepts key-value pairs:

- `type` — volume, bind, or tmpfs
- `source` (or `src`) — volume name or host path (for bind mounts)
- `target` (or `dst`) — path inside the container
- `readonly` — makes the mount read-only

Both syntaxes work; `--mount` is preferred in scripts and documentation for clarity.

6.5 Docker Volumes Cheat Sheet

Create a named volume:

```
docker volume create my-data
```

List all volumes:

```
docker volume ls
```

Inspect a volume (shows the mount point on the host):

```
docker volume inspect my-data
```

Remove a volume:

```
docker volume rm my-data
```

Remove all unused volumes:

```
docker volume prune
```

6.6 Practical Example: Nginx with Persistent Content

Let us demonstrate volumes with our Nginx container. We will create a volume, write content into it from a temporary container, then serve that content with Nginx — and show that the content survives a container replacement.

6.6.1 Create the volume

```
docker volume create web-content
```

6.6.2 Populate the volume

```
docker run --rm \
  --mount type=volume,source=web-content,target=/data \
  alpine sh -c 'echo "Content from a volume" > /data/index.html'
```

This starts a temporary Alpine container, mounts the volume at `/data`, writes a file, and immediately exits and removes itself (`--rm`). The volume and its content persist.

6.6.3 Serve the content with Nginx

```
docker run -d \  
  --name nginx-volume \  
  -p 8080:80 \  
  --mount type=volume,source=web-content,target=/usr/share/nginx/html \  
  nginx:alpine
```

Open <http://localhost:8080> — you see the page from the volume.

6.6.4 Replace the container — data survives

```
docker rm -f nginx-volume  
  
docker run -d \  
  --name nginx-volume-new \  
  -p 8080:80 \  
  --mount type=volume,source=web-content,target=/usr/share/nginx/html \  
  nginx:alpine
```

Open <http://localhost:8080> again. The same page is still there. The volume outlived the container.

6.6.5 Clean up

```
docker rm -f nginx-volume-new  
docker volume rm web-content
```

6.7 Volumes vs Bind Mounts: When to Use Which

	Volumes	Bind Mounts
Managed by Docker	Yes	No
Portable across machines	Yes (with volume drivers)	No (path must exist on host)
Best for production data	Yes	No
Best for dev code sharing	No	Yes
Access from host directly	Via <code>docker volume inspect</code>	Direct filesystem path

	Volumes	Bind Mounts
Works on Windows containers	Yes	Limited

6.8 Conclusion

In this article we covered:

- Why containers are ephemeral by default and why that matters
- The three Docker storage types: **volumes**, **bind mounts**, and **tmpfs**
- The modern **--mount** syntax vs the older **-v** flag
- Essential volume commands: `create`, `ls`, `inspect`, `rm`, `prune`
- A practical example showing that volume data survives container replacement
- A comparison table to decide when to use volumes vs bind mounts

The [next article](#) introduces **Docker Compose**: how to manage a multi-container application with a single YAML file instead of juggling multiple `docker run` commands.

If you enjoyed this article, don't forget to **give it a clap** , **share it with your friends** , and **follow me for more tips and tutorials on software development** . Your support helps me create more content like this — thank you!

7 How Docker Compose Works




Figure 10: How Docker Compose Works

7.1 Introduction

This is the sixth article of the **Getting Started with Docker** series. In the [previous article](#) we learned how to persist data with volumes. So far, to run our application we have been juggling multiple `docker run` commands — one for the network, one for each container, one for each volume. This quickly becomes hard to manage and impossible to share reliably with a team.

Docker Compose solves this: you describe your entire application in a single YAML file and launch everything with one command.

7.2 What is Docker Compose?

Docker Compose is a tool for defining and running multi-container Docker applications. You describe your services, networks, and volumes in a `docker-compose.yml` file, and Docker Compose translates that into the correct sequence of docker commands.

Important: the original `docker-compose` command (v1, written in Python) was deprecated in May 2023 and is no longer maintained. The current tool is **docker compose** (v2, written in Go), which is a plugin built directly into the Docker CLI. Always use `docker compose` (with a space, not a hyphen). If you are on Podman, use `podman compose` which is compatible with the same YAML format.

You can think of `docker compose` as a wrapper that orchestrates all the `docker` commands you would otherwise run manually — but declaratively, from a single file.

7.3 The `docker-compose.yml` File

The Compose file is a YAML document. Its top-level keys are:

- `services` — the containers that make up your application
- `volumes` — named volumes used by the services
- `networks` — custom networks (if not specified, Compose creates a default bridge network automatically)

7.3.1 A Note on the version Field

Older tutorials and documentation show a `version: '3.x'` field at the top of the Compose file. This field is **deprecated** and should be omitted. Modern Docker Compose uses the [Compose Specification](#) which does not require a version declaration.




Figure 11: Docker Compose: YAML file maps to running containers, network, and volumes

7.4 Practical Example: Frontend and Backend with Compose

We will rebuild the two-container application from the networking article using Docker Compose, adding a named volume for the frontend content.

7.4.1 Project structure

```
my-app/  
  docker-compose.yml  
  frontend/  
    index.html  
  backend/  
    index.html
```

Create the directories and content:

```
mkdir -p my-app/frontend my-app/backend  
cd my-app  
  
cat > frontend/index.html  
  
Frontend  
  
Frontend  
This app is managed by Docker Compose.  
  
EOF  
  
cat > backend/index.html << 'EOF'  
{ "status": "ok", "message": "Hello from the backend!" }  
EOF
```

7.4.2 The docker-compose.yml file

```
services:  
  frontend:  
    image: nginx:alpine  
    container_name: frontend  
    ports:  
      - "8080:80"  
    volumes:
```

```

    - ./frontend:/usr/share/nginx/html:ro
networks:
  - app-network
depends_on:
  - backend

backend:
  image: nginx:alpine
  container_name: backend
  volumes:
    - ./backend:/usr/share/nginx/html:ro
  networks:
    - app-network

networks:
  app-network:
    driver: bridge

```

Key things to notice:

- **No version: field** — not needed with modern Compose
- **depends_on** — ensures the backend container starts before the frontend. For health-based ordering, you can use `condition: service_healthy` together with a healthcheck definition
- **Networks are created automatically** — Compose creates `app-network` for you
- **Bind mounts** use relative paths from the directory containing `docker-compose.yml`

7.5 Docker Compose Commands

The most important commands:

```

# Start all services (build images if needed, create networks and volumes)
docker compose up -d

# Stop and remove containers, networks (volumes are preserved)
docker compose down

```

```
# Stop and remove everything including volumes
docker compose down -v

# View logs of all services
docker compose logs

# Follow logs in real time
docker compose logs -f

# List running containers managed by this Compose file
docker compose ps

# Execute a command inside a running service container
docker compose exec frontend sh

# Rebuild images (useful after changing a Dockerfile)
docker compose build

# Pull the latest versions of all images
docker compose pull

# Restart a single service
docker compose restart frontend
```

The key difference from docker commands is that all of these act on the **entire application** defined in the Compose file, not a single container.

7.6 Environment Variables

Hard-coded values in `docker-compose.yml` are fine for development, but for anything sensitive (passwords, API keys) or environment-specific (URLs, port numbers) you should use variables.

Docker Compose automatically reads a file named `.env` in the same directory:

```
# .env
NGINX_PORT=8080
APP_ENV=development
```

Then reference them in `docker-compose.yml`:

```
services:
  frontend:
    ports:
      - "${NGINX_PORT}:80"
    environment:
      - APP_ENV=${APP_ENV}
```

Never commit `.env` files containing secrets to version control. Add `.env` to your `.gitignore`.

7.7 Starting the Application

From the `my-app/` directory:

```
docker compose up -d
```

Compose will:

1. Create the `app-network` bridge network
2. Start the backend container
3. Start the frontend container
4. Map port 8080 on your host to port 80 on the frontend container

Open `http://localhost:8080` — the frontend is running. Check that both services are up:

```
docker compose ps
```

Verify the backend is reachable by name from the frontend:

```
docker compose exec frontend ping -c 3 backend
```

7.7.1 Clean up

```
docker compose down
```

7.8 Conclusion

In this article we covered:

- What Docker Compose is and why it replaces manual `docker run` scripts
- The difference between `docker-compose v1` (deprecated) and **`docker compose v2`** (current)

- The structure of a `docker-compose.yml` file: services, volumes, networks
- Why the `version:` field is no longer needed
- Essential Compose commands: `up`, `down`, `logs`, `ps`, `exec`, `build`
- How to manage configuration with **environment variables** and `.env` files

The **next article** closes the series with **Docker security best practices**: how to run containers safely, limit their privileges, and scan images for vulnerabilities.

If you enjoyed this article, don't forget to **give it a clap** , **share it with your friends** , and **follow me for more tips and tutorials on software development** . Your support helps me create more content like this — thank you!

8 Docker Security Best Practices



Figure 12: Docker Security Best Practices

8.1 Introduction

This is the seventh and final article of the **Getting Started with Docker** series. In the [previous article](#) we orchestrated a multi-container application with Docker Compose. Now we turn our attention to **security**.

Containers are isolated — but isolation is not the same as security. A container running as root, with a writable filesystem and full Linux capabilities, is a significant risk if it is ever compromised. The good news is that Docker provides several mechanisms to reduce the attack surface, and most of them require only a few lines of configuration.

8.2 1. Never Run as Root

By default, processes inside a container run as **root** (UID 0). This means that if an attacker exploits a vulnerability in your application, they have root access inside the container — and depending on how the container is configured, they may be able to escape to the host.

The fix is simple: add a `USER` instruction to your Dockerfile.

```
FROM nginx:alpine

COPY html /usr/share/nginx/html
```



Figure 13: Docker security best practices overview

```
# Switch to the non-root user that the nginx image already provides
USER nginx

EXPOSE 80

CMD ["nginx", "-g", "daemon off;"]
```

The official `nginx:alpine` image already includes a non-root `nginx` user. For your own images, create a dedicated user:

```
FROM alpine:3.20

RUN addgroup -S appgroup && adduser -S appuser -G appgroup

COPY --chown=appuser:appgroup app /app

USER appuser

CMD ["/app/start.sh"]
```

The `--chown` flag in `COPY` ensures the files are owned by the correct user from the start.

You can also enforce non-root at runtime without changing the Dockerfile:

```
docker run --user 1001:1001 nginx:alpine
```

8.3 2. Use a Read-Only Filesystem

If your application does not need to write to its own filesystem, make it read-only:

```
docker run --read-only -p 8080:80 nginx:alpine
```

If the application needs to write to specific paths (temp files, logs), mount those paths as tmpfs:

```
docker run --read-only \  
  --tmpfs /tmp \  
  --tmpfs /var/cache/nginx \  
  --tmpfs /var/run \  
  -p 8080:80 \  
  nginx:alpine
```

A read-only filesystem means that even if an attacker gains code execution inside the container, they cannot modify the application binaries or install tools.

In Docker Compose:

```
services:  
  frontend:  
    image: nginx:alpine  
    read_only: true  
    tmpfs:  
      - /tmp  
      - /var/cache/nginx  
      - /var/run
```

8.4 3. Drop Linux Capabilities

Linux capabilities are fine-grained divisions of root privileges. By default, Docker grants a container a subset of them. You can drop all capabilities and add back only what your application actually needs:

```
docker run --cap-drop=ALL --cap-add=NET_BIND_SERVICE nginx:alpine
```

NET_BIND_SERVICE allows binding to ports below 1024 (like port 80). Most application containers do not need any other capability.

In Docker Compose:

```
services:
  frontend:
    image: nginx:alpine
    cap_drop:
      - ALL
    cap_add:
      - NET_BIND_SERVICE
```

8.5 4. Never Put Secrets in Environment Variables (or Images)

A common mistake is passing passwords or API keys through environment variables in `docker-compose.yml` or baking them into the image with `ENV`. Both approaches expose secrets in `docker inspect` output and in image layers.

What to do instead:

- Use **Docker secrets** (available in Swarm mode) which mount secrets as files in `/run/secrets/`
- Use a **secrets manager** (HashiCorp Vault, AWS Secrets Manager, etc.)
- As a minimum, use a `.env` file that is **never committed to version control** and reference it in Compose

```
# docker-compose.yml
services:
  app:
    image: my-app
    env_file:
      - .env.secrets  # never committed to git
```

Add to `.gitignore`:

```
.env
.env.secrets
*.env
```

8.6 5. Use Minimal Base Images

Every package in your base image is a potential attack surface. Prefer minimal base images:

- `alpine` — ~5MB, very popular, uses `musl libc`
- `distroless` (Google) — no shell, no package manager, only the runtime and your app
- `scratch` — completely empty, for statically compiled binaries (Go, Rust)

For our Nginx example we already use `nginx:alpine`. Going even smaller would mean building a custom Nginx binary, which is rarely necessary.

The rule: **if it is not needed, it should not be in the image.**

8.7 6. Scan Images for Known Vulnerabilities

Even minimal images can contain packages with known CVEs. Docker provides a built-in scanner called **Docker Scout**:

```
# Quick vulnerability overview
docker scout quickview nginx:alpine

# Full CVE list
docker scout cves nginx:alpine

# Recommendations to fix vulnerabilities
docker scout recommendations nginx:alpine
```

Docker Scout is available in Docker Desktop and as a CLI plugin. It integrates with Docker Hub and can be added to CI/CD pipelines to fail a build when high-severity vulnerabilities are detected.

An alternative open-source scanner is **Trivy**:

```
trivy image nginx:alpine
```

Make image scanning part of your build pipeline, not an afterthought.

8.8 7. Keep Images Up to Date

A scanned image with no vulnerabilities today may have CVEs tomorrow. Pin your base image to a specific digest rather than a floating tag, and automate rebuild checks:

```
# Pin to a specific digest (reproducible, auditable)
FROM nginx:alpine@sha256:...
```

Many teams use tools like **Renovate** or **Dependabot** to automatically open PRs when a newer

base image is available.

8.9 Putting It All Together

Here is a Dockerfile for our Nginx image that applies all the practices above:

```
FROM nginx:alpine

# Copy content
COPY --chown=nginx:nginx html /usr/share/nginx/html

# Drop to non-root
USER nginx

EXPOSE 80

CMD ["nginx", "-g", "daemon off;"]
```

And the corresponding docker-compose.yml:

```
services:
  frontend:
    build: .
    image: my-nginx:secure
    container_name: frontend
    ports:
      - "8080:80"
    read_only: true
    tmpfs:
      - /tmp
      - /var/cache/nginx
      - /var/run
    cap_drop:
      - ALL
    cap_add:
      - NET_BIND_SERVICE
    networks:
      - app-network
```

```
networks:
  app-network:
    driver: bridge
```

8.10 Conclusion

In this article we covered:

- Running containers as a **non-root user** with USER in the Dockerfile
- Using a **read-only filesystem** with `--read-only` and `tmpfs` for writable paths
- **Dropping Linux capabilities** with `--cap-drop=ALL` and adding back only what is needed
- Keeping **secrets out of images and environment variables**
- Choosing **minimal base images** to reduce attack surface
- **Scanning images** for vulnerabilities with Docker Scout or Trivy
- **Keeping images up to date** with pinned digests and automated updates

This concludes the **Getting Started with Docker** series. You now have a solid foundation to build, network, persist, orchestrate, and secure containerized applications. The natural next step is **Docker Swarm** for multi-host deployments, or **Kubernetes** for large-scale container orchestration.

If you enjoyed this article, don't forget to **give it a clap** , **share it with your friends** , and **follow me for more tips and tutorials on software development** . Your support helps me create more content like this — thank you!

9 Glossary

Bind Mount A Docker storage mechanism that maps a specific file or directory on the host filesystem directly into a container. Unlike named volumes, bind mounts are not managed by Docker and depend on the host's directory structure.

BuildKit The modern build engine for Docker, enabled by default since Docker 23.0. BuildKit provides faster builds, better caching, and new features such as cache mounts and secret mounts in Dockerfiles.

cgroups (Control Groups) A Linux kernel feature that limits, accounts for, and isolates the resource usage (CPU, memory, disk I/O, network) of a collection of processes. Docker uses cgroups to enforce per-container resource limits.

Container A runnable instance of a Docker image. Containers are isolated from each other and from the host using Linux namespaces and cgroups, but share the host OS kernel.

containerd The industry-standard container runtime that Docker Engine uses under the hood since Docker 1.11. containerd manages the complete container lifecycle: image transfer, storage, execution, and supervision.

Docker Compose A tool for defining and running multi-container Docker applications using a declarative YAML file (`docker-compose.yml`). The current version is `docker compose` (v2, a Go plugin); the older `docker-compose` (v1, Python) was deprecated in May 2023.

Docker Desktop A GUI application for Mac and Windows that bundles Docker Engine, Docker CLI, Docker Compose, and Kubernetes. Requires a paid subscription for enterprise use.

Docker Hub The default public registry for Docker images. Hosts official images maintained by Docker and thousands of community images.

Docker Scout A vulnerability scanning tool integrated into the Docker CLI and Docker Desktop. Analyses images for known CVEs and provides remediation recommendations.

Dockerfile A plain-text file containing a sequence of instructions that tell Docker how to build an image. Each instruction creates a new layer in the image.

Image A read-only, layered template used to create containers. Images are built from Dockerfiles and stored in registries. Each layer represents a filesystem change.

Layer Cache The mechanism by which Docker reuses unchanged layers from previous builds. Layers are rebuilt only when an instruction or its inputs change, making subsequent builds significantly faster.

Namespace A Linux kernel feature that provides isolation for system resources. Docker uses multiple namespaces per container: PID (processes), NET (network interfaces), MNT (filesystem mounts), UTS (hostname), IPC, and USER.

OCI (Open Container Initiative) A set of open industry standards for container formats and runtimes. Both Docker and Podman produce OCI-compliant images, which means images built with one tool can be run by the other.

Podman A daemonless, rootless container engine that is fully compatible with the Docker CLI. An open-source alternative to Docker Desktop, particularly useful on corporate machines where Docker Desktop requires a commercial licence.

Registry A service that stores and distributes Docker images. Docker Hub is the default public registry. Private registries include AWS ECR, GitHub Container Registry, and self-hosted solutions.

Reverse Proxy A server that sits in front of backend services and forwards client requests to them. Nginx is commonly used as a reverse proxy in containerised applications.

tmpfs Mount A Docker storage mechanism that stores data in the host's RAM. The data is never written to disk and is lost when the container stops. Useful for sensitive data and high-speed temporary storage.

Volume A Docker-managed storage unit that persists data outside the container's filesystem. Volumes survive container removal and can be shared across multiple containers.

WSL2 (Windows Subsystem for Linux 2) A compatibility layer in Windows 10/11 that runs a real Linux kernel inside a lightweight VM. Docker Desktop uses WSL2 as its Linux backend on Windows.

10 What's Next?

You have reached the end of this book. If you have worked through each chapter, you now know how to:

- Pull and run containers from Docker Hub
- Explain the difference between containers and virtual machines at a kernel level
- Write Dockerfiles, build custom images, and understand the layer cache
- Create user-defined bridge networks and let containers resolve each other by name
- Persist data with named volumes and bind mounts
- Orchestrate a multi-container application with Docker Compose
- Harden a container image with security best practices

That is a complete, practical foundation. Most day-to-day Docker work does not require anything beyond what you have learned here.

10.1 Where to Go From Here

10.1.1 Docker Swarm

Docker Swarm is Docker's built-in clustering and orchestration mode. It lets you deploy multi-container applications across multiple hosts with a single `docker stack deploy` command. If you need high availability without the complexity of Kubernetes, Swarm is the right next step.

10.1.2 Kubernetes

Kubernetes (K8s) is the industry-standard container orchestrator for large-scale production deployments. It builds on the same container concepts you now know and adds automated roll-outs, self-healing, horizontal scaling, and much more. The [Getting Started with Kubernetes](#) series on Code4Projects is a natural continuation.

10.1.3 CI/CD with Docker and GitHub Actions

Automating the build, test, and push of your images in a CI/CD pipeline is the next productivity leap. GitHub Actions has first-class Docker support and can build multi-platform images (arm64 + amd64) with a single workflow file.

10.1.4 Docker Security (deeper)

This book covered the essentials. For production workloads, the next steps are: image signing with Docker Content Trust, runtime security with Falco, network policies, and secrets manage-

ment with HashiCorp Vault or AWS Secrets Manager.

10.2 Stay in Touch

If this book helped you, consider:

- Subscribing to the **Code4Projects newsletter** on Substack for new articles and future ebook releases
- Starring the GitHub repository at github.com/sasadangelo/docker-tutorials
- Sharing the book with a colleague who is just getting started with containers

Thank you for reading.

— *Salvatore D'Angelo*

code4projects.com

*The best way to learn Docker is to use it every day.
Start with one container. Build from there.*